

Data 188      Introduction to Deep Learning  
 Spring 2026    Eric Kim      Discussion 01

Welcome to Data 188 - we're excited to have you here and have some *deep* conversations with you!  
 This discussion will cover some matrices, vectors, and gradients review.

## 1. Matrix Multiplication

Compute each of the following matrix multiplications by hand. If the multiplication is not valid, explain why.

$$(a) \begin{bmatrix} 1 & 6 \\ 7 & 5 \end{bmatrix} \begin{bmatrix} 3 \\ 0 \\ 2 \end{bmatrix}$$

**Solution:** The first matrix is  $2 \times 2$  and the second is  $3 \times 1$ . The number of columns in the first matrix (2) does not match the number of rows in the second matrix (3), so this multiplication is invalid.

$$(b) \begin{bmatrix} 4 & 3 & 6 \\ 2 & 1 & 0 \\ 0 & 5 & 7 \end{bmatrix} \begin{bmatrix} 2 \\ 1 \\ 3 \end{bmatrix}$$

**Solution:**

$$\begin{bmatrix} 4 \cdot 2 + 3 \cdot 1 + 6 \cdot 3 \\ 2 \cdot 2 + 1 \cdot 1 + 0 \cdot 3 \\ 0 \cdot 2 + 5 \cdot 1 + 7 \cdot 3 \end{bmatrix} = \begin{bmatrix} 8 + 3 + 18 \\ 4 + 1 + 0 \\ 0 + 5 + 21 \end{bmatrix} = \begin{bmatrix} 29 \\ 5 \\ 26 \end{bmatrix}$$

$$(c) \begin{bmatrix} 0 & 6 & 5 \\ 2 & 3 & 3 \\ 5 & 7 & 8 \end{bmatrix} \begin{bmatrix} 1 & 9 \\ 1 & 4 \\ 3 & 0 \end{bmatrix}$$

**Solution:**

$$\begin{bmatrix} 0 \cdot 1 + 6 \cdot 1 + 5 \cdot 3 & 0 \cdot 9 + 6 \cdot 4 + 5 \cdot 0 \\ 2 \cdot 1 + 3 \cdot 1 + 3 \cdot 3 & 2 \cdot 9 + 3 \cdot 4 + 3 \cdot 0 \\ 5 \cdot 1 + 7 \cdot 1 + 8 \cdot 3 & 5 \cdot 9 + 7 \cdot 4 + 8 \cdot 0 \end{bmatrix} = \begin{bmatrix} 0 + 6 + 15 & 0 + 24 + 0 \\ 2 + 3 + 9 & 18 + 12 + 0 \\ 5 + 7 + 24 & 45 + 28 + 0 \end{bmatrix} = \begin{bmatrix} 21 & 24 \\ 14 & 30 \\ 36 & 73 \end{bmatrix}$$

## 2. Jacobians

Recall that the Jacobian (or Jacobian matrix) of a vector-valued function is the matrix of all its first-order partial derivatives. If  $\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^k$  is a vector-valued function with  $n$  inputs and  $k$  outputs, then its *Jacobian*

or *gradient*  $\nabla \mathbf{f}$  is an  $k \times n$  matrix defined as

$$\nabla \mathbf{f} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \frac{\partial f_1}{\partial x_1} & \dots & \frac{\partial f_n}{\partial x_1} \\ \frac{\partial f_1}{\partial x_2} & \frac{\partial f_2}{\partial x_2} & \dots & \frac{\partial f_n}{\partial x_2} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial f_1}{\partial x_k} & \frac{\partial f_2}{\partial x_k} & \dots & \frac{\partial f_n}{\partial x_k} \end{bmatrix}$$

**Numerator convention:** Note that in this class, we will use the "numerator" convention rather than the "denominator" convention for the dimensions of the Jacobian. In other courses or textbooks, you may see that if  $\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^k$  then  $\nabla \mathbf{f} \in \mathbb{R}^{n \times k}$ . However, in this course,  $\nabla \mathbf{f} \in \mathbb{R}^{k \times n}$ .

We call it the "numerator" convention because if we call the input vector  $\mathbf{x} \in \mathbb{R}^n$  and the output vector  $\mathbf{y} \in \mathbb{R}^k$ , then each entry of the Jacobian is of the form  $\frac{\partial y_i}{\partial x_j}$ , where the variable in the numerator corresponds to the output dimension and the variable in the denominator corresponds to the input dimension. In other words, the Jacobian dimensions are  $\text{dim}(\text{numerator}) \times \text{dim}(\text{denominator})$ .

**Matrix-valued functions:** Also note that in this class, we will often work with functions that accept matrices (not just vectors) as input/output. When we define the Jacobian of these matrix-valued functions, the "true" shape of the Jacobian is still 2D, not 3D/4D.

Example: Suppose  $\mathbf{f}$  takes an  $n \times k$  matrix, and outputs a scalar. Its Jacobian  $\nabla \mathbf{f}(\mathbf{x})$  has shape  $1 \times n * k$ , where the first dimension is "flattened" (eg row-wise). Notably, the shape of  $\nabla \mathbf{f}(\mathbf{x})$  is NOT  $n \times k$ .

However, for convenience, people often represent  $\nabla \mathbf{f}(\mathbf{x})$  as a matrix (with shape  $n \times k$ ) rather than as a long vector (with shape  $1 \times n * k$ ). This is fine, as long as you recognize when this happens!

**Instructions:** For each of the following functions, state the dimensions of the Jacobian and compute it.

$$(a) \mathbf{f}(x, y) = \begin{bmatrix} x^3 y \\ \cos^2 y + 5x \\ ye^x \end{bmatrix}$$

**Solution:** The input is a 2-dimensional vector  $(x, y)$  and the output is a 3-dimensional vector, so the Jacobian will be of size  $3 \times 2$ .

$$\begin{aligned} \nabla \mathbf{f} &= \begin{bmatrix} \frac{\partial f_1}{\partial x} & \frac{\partial f_1}{\partial y} \\ \frac{\partial f_2}{\partial x} & \frac{\partial f_2}{\partial y} \\ \frac{\partial f_3}{\partial x} & \frac{\partial f_3}{\partial y} \end{bmatrix} \\ &= \begin{bmatrix} 3x^2 y & x^3 \\ 5 & -2 \sin y \cos y \\ ye^x & e^x \end{bmatrix} \end{aligned}$$

$$(b) \mathbf{f}\left(\begin{bmatrix} p & q \\ r & s \end{bmatrix}\right) = \begin{bmatrix} p + 4q \\ \sin(r^2s) \end{bmatrix}$$

**Solution:** The input is a  $2 \times 2$  matrix and the output is a 2-dimensional vector, so the Jacobian will be of size  $2 \times 4$ .

To compute the Jacobian, we flatten the input matrix in row-major order:

$$\begin{bmatrix} p \\ q \\ r \\ s \end{bmatrix}$$

and then compute the partial derivatives like before:

$$\begin{aligned} \nabla \mathbf{f} &= \begin{bmatrix} \frac{\partial f_1}{\partial p} & \frac{\partial f_1}{\partial q} & \frac{\partial f_1}{\partial r} & \frac{\partial f_1}{\partial s} \\ \frac{\partial f_2}{\partial p} & \frac{\partial f_2}{\partial q} & \frac{\partial f_2}{\partial r} & \frac{\partial f_2}{\partial s} \end{bmatrix} \\ &= \begin{bmatrix} 1 & 4 & 0 & 0 \\ 0 & 0 & 2rs \cos(r^2s) & r^2 \cos(r^2s) \end{bmatrix} \end{aligned}$$

### 3. Deriving the Linear Softmax Gradient Update Equations

Recall from [Lecture 2](#) that we can formulate linear softmax regression as an optimization problem where we want to find the *parameters* (aka *weights*)  $\theta$  that minimizes the cross-entropy loss function  $f(\theta)$  using the gradient descent algorithm. In gradient descent, we iteratively update  $\theta$  using the *gradient update equation*:

$$\theta := \theta - \alpha \nabla_{\theta} f(\theta)$$

where  $\alpha$  is the *learning rate* (aka *step size*).

In this problem, we will derive the gradient update equation for linear softmax regression step-by-step.

Consider a  $k$ -class classification problem where we have:

- Training data:  $x^{(i)} \in \mathbb{R}^n$ ,  $y^{(i)} \in \{1, \dots, k\}$  for  $i = 1, \dots, m$
- $n$  = dimensionality of input data
- $k$  = number of classes/labels
- $m$  = number of points in the training set

Our hypothesis function maps inputs  $x \in \mathbb{R}^n$  to  $k$ -dimensional vectors  $h : \mathbb{R}^n \rightarrow \mathbb{R}^k$  where  $h(x)$  indicates some measure of "belief" in how much likely the label is to be class  $i$  (i.e., "most likely" prediction is coordinate  $i$  with largest  $h(x)$ ).

A linear hypothesis function uses a linear operator (i.e. matrix multiplication) for this transformation  $h_{\theta}(x) = \theta^T x$  for parameters  $\theta \in \mathbb{R}^{n \times k}$ .

To convert the values from our hypothesis function to probabilities (e.g. make sure they are all non-negative values that sum to 1), we apply the softmax function  $z$  to  $h_{\theta}(x)$ :

$$z_i = p(\text{label} = i) = \frac{e^{h_i(x)}}{\sum_{j=1}^k e^{h_j(x)}}$$

Then the *cross-entropy loss* function  $l_{ce}$  is defined as:

$$\begin{aligned} l_{ce}(h(x), y) &= -\log p(\text{label} = y) \\ &= -\log z_y \\ &= -\log \frac{e^{h_y(x)}}{\sum_{j=1}^k e^{h_j(x)}} \\ &= -(\log e^{h_y(x)} - \log \sum_{j=1}^k e^{h_j(x)}) \\ &= -h_y(x) + \log \sum_{j=1}^k e^{h_j(x)} \\ &= f(\theta) \end{aligned}$$

Recall the multivariate chain rule for computing the derivative of a composition of functions. Let  $h = \theta^T x$ . We can compute  $\nabla_\theta f(\theta)$  by breaking it down into the product of 2 partial derivatives:

$$\begin{aligned} \nabla_\theta f(\theta) &= \frac{\partial l_{ce}(h, y)}{\partial \theta} \\ &= \frac{\partial l_{ce}(h, y)}{\partial h} \cdot \frac{\partial h}{\partial \theta} \end{aligned}$$

- (a) Show that  $\frac{\partial l_{ce}(h, y)}{\partial h} = (z - e_y)^T$  where  $e_y$  is the one-hot vector with a 1 in coordinate  $y$  and 0s elsewhere.

**Hint 1:** Start by computing  $\frac{\partial l_{ce}(h, y)}{\partial h_i}$  and then generalize that into vector form.

**Hint 2:** How does the expression in Hint 1 differ when  $i = y$  versus when  $i \neq y$ ?

**Solution:** Following the hint:

$$\begin{aligned} \frac{\partial l_{ce}(h, y)}{\partial h_i} &= \frac{\partial}{\partial h_i} (-h_y + \log \sum_{j=1}^k e^{h_j}) \\ &= -1\{i = y\} + \frac{e^{h_i}}{\sum_{j=1}^k e^{h_j}} \end{aligned}$$

where  $-1\{i = y\}$  is an indicator function that is  $-1$  if  $i = y$  and 0 otherwise.

Since  $\frac{e^{h_i}}{\sum_{j=1}^k e^{h_j}} = z_i$  and addition is commutative, we can rewrite the above as:

$$\frac{\partial l_{ce}(h, y)}{\partial h_i} = z_i - 1\{i = y\}$$

Therefore, in vector form, we have:

$$\frac{\partial l_{ce}(h, y)}{\partial h} = \frac{\partial l_{ce}(h, y)}{\partial h} = (z - e_y)^T$$

(We need to transpose to get the correct shape based on numerator convention.)

See also: Lecture 2, slide 28.

- (b) Compute  $\frac{\partial h}{\partial \theta}$ .

**Solution:** Recall that  $h = \theta^T x$ , so  $h_i$  only depends on the  $i$ -th row of  $\theta$ . Thus,  $\frac{\partial h_i}{\partial \theta}$  is a  $n \times k$  matrix where all columns are 0 except for the  $i$ -th column, which is equal to  $x$ . (Technically it is a  $1 \times nk$  matrix but we can reshape it to  $n \times k$  for convenience.)

See also: Lecture 2, slide 29.

- (c) Recall that in practice, we do not want to compute  $\frac{\partial h}{\partial \theta}$  directly because it would take too much compute time and memory. Instead, we can use the results from parts (a) and (b) and take advantage of structure to compute  $\nabla_\theta f(\theta)$  more efficiently.

**Hint 1:** A matrix-vector multiplication  $Ax$  can be expressed as scaling the  $i$ -th column of  $A$  by the  $i$ -th entry in  $x$  and summing each scaled column. That is,  $Ax = \sum_{i=1}^n A_{:,i}x_i$  where  $A_{:,i}$  is the  $i$ -th column of  $A$ .

**Hint 2:** It can be useful to define a vector  $c = \frac{\partial l_{ce}}{\partial h}$  and let  $c_i$  be its  $i$ -th entry.

**Hint 3:** A matrix of the form

$$\begin{bmatrix} \vdots & \vdots & \vdots & \vdots \\ v_1 u & v_2 u & \cdots & v_k u \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix}$$

where  $u$  and  $v$  are column vectors can be expressed as  $u * v^T$ .

**Solution:**

$$\begin{aligned} \frac{\partial l_{ce}(h, y)}{\partial h} \cdot \frac{\partial h}{\partial \theta} &= \sum_{i=1}^k \frac{\partial l_{ce}(h, y)}{\partial h_i} \cdot \frac{\partial h_i}{\partial \theta} \\ &= \begin{bmatrix} \vdots & \vdots & \vdots & \vdots \\ c_1 x & 0 & \cdots & 0 \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix} + \begin{bmatrix} \vdots & \vdots & \vdots & \vdots \\ 0 & c_2 x & \cdots & 0 \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix} + \cdots + \begin{bmatrix} \vdots & \vdots & \vdots & \vdots \\ 0 & 0 & \cdots & c_k x \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix} \\ &= \begin{bmatrix} \vdots & \vdots & \vdots & \vdots \\ c_1 x & c_2 x & \cdots & c_k x \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix} \\ &= x * \frac{\partial l_{ce}}{\partial h} \end{aligned}$$

$$= x(z - e_y)^T$$

**Contributors:**

- Rebecca Dang.